

CS495 Capstone: Final Report

Title: Integer Optimization for Logistics Assignment, with a Three-Language Branch & Bound Study **Author:** Peter Kamau **Instructor:** Prof. Dr. Pedro Albuquerque **Course:** CS495 Capstone in Data Science, Bellevue College **Date:** May 2026

Abstract

This capstone studies discrete optimization in a logistics setting and extends the inquiry into how the same algorithm performs across implementation languages. The project has two tracks:

1. **Production track.** A driver-to-region assignment binary ILP, modeled with PuLP and solved with COIN-CBC.
2. **Research track.** A hand-built Branch & Bound (B&B) 0-1 knapsack solver implemented in three languages (Python, C++, x86-64 assembly), benchmarked against brute-force enumeration.

What ties the two tracks together is that the ILP solver behind the production track (CBC) is a Branch & Bound implementation. Studying B&B directly therefore informs the engine the production model depends on.

1. Motivation

Driver-to-region assignment is a canonical assignment problem: binary choices, linear capacity/coverage constraints, a cost-minimizing objective. Logistics operators encounter variants daily, and bad assignments produce uneven workloads, missed coverage, and avoidable cost. Integer linear programming is a well-understood method for these problems and is solved in practice by Branch & Bound (with cuts and heuristics).

To understand how that engine behaves, this project also implements B&B by hand on a related, smaller problem (0-1 knapsack) where the optimum can be hand-verified. The same algorithm is implemented in three languages so language overhead can be cleanly separated from algorithmic overhead.

2. Production Track: Driver-to-Region ILP

2.1 Formulation

Let $x_{ij} \in \{0, 1\}$ be 1 if driver i is assigned to region j . Given costs c_{ij} and an eligibility flag e_{ij} :

$$\begin{array}{llll} \text{minimize} & \sum c_{ij} \cdot x_{ij} & & (\text{over eligible pairs only}) \\ \text{subject to} & \sum_j x_{ij} \leq 1 & & (\text{each driver covers at most one region}) \\ & \sum_i x_{ij} \geq 1 & & (\text{each region receives at least one driver}) \\ & x_{ij} \in \{0, 1\} & & \end{array}$$

Implementation in `src/optimization_model.py`. Sample data in `data/sample_driver_region_data.csv`: 4 drivers, 3 regions, eligibility-filtered cost matrix.

2.2 Result

On the sample instance:

- **Optimal cost:** 13
- **Assignments:** D2 $\rightarrow$ R2 (cost 4), D3 $\rightarrow$ R3 (cost 5), D4 $\rightarrow$ R1 (cost 4)
- **Coverage:** 3/3 regions covered with 3 of 4 drivers
- **Solver status:** Optimal solution found (PuLP / COIN-CBC, default settings)

D1 is unassigned. Every D1 pair was more expensive than the alternatives selected for those regions. All constraints satisfied.

3. Research Track: Branch & Bound in Three Languages

3.1 Algorithm

Items are pre-sorted by `value/weight` ratio descending. DFS over the binary include/exclude tree, ordered so include is tried first (the greedy choice). At each node, an upper bound is computed by the LP relaxation (fractional knapsack). If the bound $\leq$ the current incumbent, the subtree is pruned.

Implementations:

- Python: `src/knapsack_branch_bound.py`
- C++: `src/knapsack_branch_bound.cpp`
- x86-64 NASM (Win64 ABI, integer-floor bound): `src/knapsack_branch_bound.asm + src/knapsack_bb_asm_wrapper.c`

Brute-force baselines (`src/knapsack_brute_force.{py,cpp,asm}`) enumerate all 2^n subsets and were run on the same instances for comparison.

[image]

Figure 0. Branch & Bound on a 4-item instance (capacity = 8). Of the 16 possible subsets, the algorithm visits only 9 nodes: 3 pruned by the LP-relaxation bound, 2 pruned as infeasible. The optimal subset $\{B, C\}$ with value 22 is found early and powers the rest of the pruning. At $n = 25$ the analogous reduction is roughly one million to one.

3.2 Benchmark Protocol

All benchmarks were run on a single Windows 10 desktop using the `perf_counter_ns` high-resolution timer (Python `time.perf_counter`) and `std::chrono::steady_clock` (C++/ASM). Each instance was timed once per language per algorithm; the times reported below are single-run wall-clock measurements. At instance sizes below $n \approx 10$ the absolute times approach the timer's $\sim 1 \mu\text{s}$ resolution floor and ratios computed across those data points should be read as noise-dominated. From $n \geq 10$ the times are well above the floor and the relative comparisons are stable across reruns within a few percent.

Instances are the canonical PLAN-v1 example ($n = 3$) plus synthetic instances `notebooks/instance_n{10,15,20,22,25}.txt` generated with fixed random seeds so every reported time is reproducible from the same input file.

3.3 Benchmark Results

Single-run wall-clock times on six instances ($n = 3$ to 25):

n	Optimum	Brute Python	Brute C++ +	B&B Python	B&B C++ +	B&B ASM	B&B/Brute (Python)
10	490	1.26 ms	0.36 ms	20.9 μ s	6 μ s	0.5 μ s	$\sim 60\times$
15	489	60.2 ms	13.6 ms	55.5 μ s	3 μ s	0.9 μ s	$\sim 1,085\times$
20	880	2.74 s	0.46 s	94.6 μ s	5 μ s	1.8 μ s	$\sim 28,960\times$
22	914	11.79 s	1.90 s	57.6 μ s	4 μ s	1.0 μ s	$\sim 204,627\times$
25	1153	102.22 s	15.60 s	96.5 μs	5 μ s	1.7 μ s	$\sim 1,059,000\times$

Raw data: `benchmark_results_all_three.json`, `benchmark_results_bb.json`.

[image]

Figure 1. Brute force: three languages, six instances. Exponential growth in all three; the language ranking is consistent ($ASM < C++ < Python$) but does not change the curve.

[image]

Figure 2. Branch & Bound: same instances, now measured in microseconds. The y-axis units shift from seconds (Figure 1) to microseconds (Figure 2). That gap is the algorithmic effect.

3.4 What the Numbers Show

- **Algorithmic effect.** B&B closes the gap with brute force by orders of magnitude that grow with n . At $n = 25$, Python B&B is approximately one million times faster than Python brute force and returns the identical optimum.
- **Language effect.** Holding the algorithm fixed, the ranking is consistent: $ASM > C++ > Python$. Within B&B, C++ averages $\sim 13\times$ over Python and ASM averages $\sim 52\times$ over Python ($\sim 3\times$ over C++).
- **Independence.** Algorithmic and language wins multiply. The language ranking is the same regardless of which algorithm is used.

3.5 Scaling Analysis

Three quantitative checks confirm that the empirical curves match the expected $O(2^n)$ complexity and that the language gap behaves as predicted.

[image]

Figure 3. Scaling diagnostics. Left: log-linear fits of $T = a \cdot 2^n$ to each language. Middle: empirical growth factor $T(n+1) / T(n)$. Right: the Python/C++ and Python/ASM speedup ratios as functions of n .

Curve fit (left panel). Taking $T = a \cdot 2^n$ and $\log T = \log a + n \cdot \log 2$, fitting $\log T$ against n linearly yields:

Language Constant a (sec) Implied growth slope

Python	5.38×10^{-7}	$\log 2 \approx 0.301$
C++	8.56×10^{-8}	$\log 2 \approx 0.301$
ASM	9.96×10^{-8}	$\log 2 \approx 0.301$

The fits are visually indistinguishable from the data on the log axis. The slope is identical across languages. Only the constant a (per-operation cost) changes. This is the empirical confirmation that the algorithm's asymptotic class is independent of implementation language.

Empirical growth factor (middle panel). With instance steps of $\Delta n = 2, 3, 5, 7$, the expected ratio between consecutive runtimes is $2^{\Delta n}$. Selected measured ratios:

n step	Expected $2^{\Delta n}$	Python	C++	ASM
20 $\rightarrow$ 22 ($\Delta n=2$)	4.00	4.30	4.11	4.26
22 $\rightarrow$ 25 ($\Delta n=3$)	8.00	8.67	8.21	8.74

All three languages track the theoretical doubling line once $n \geq 10$. At small n (notably the $n = 3 \rightarrow 10$ step) the ASM growth factor appears inflated because the $n=3$ absolute time ($1 \mu s$) is at the timer's resolution floor. The denominator is artificially small. This is a measurement artifact, not an algorithmic difference.

Performance gap evolution (right panel). If both languages obeyed $T = a \cdot 2^n$ exactly, the speedup ratio $T_{py} / T_{compiled}$ would be constant in n . It is not. Both ratios climb with n :

n	Python / C++	Python / ASM
3	0.48 (C++ slower than Python)	4.80
10	3.55	14.15
20	5.93	20.91
25	6.55	20.95

Real measured time is $\text{overhead} + a \cdot 2^n$, not pure exponential. At small n , overhead dominates and the ratios reflect fixed costs (program startup, allocator, file I/O). As n grows, the 2^n term takes over and the ratios converge to the true per-operation cost difference between languages. The Python/ASM ratio has plateaued around $21\times$ by $n = 25$ (ASM has reached its asymptote); the Python/C++ ratio is still climbing slowly and may settle near $7\text{--}8\times$ at larger n .

4. Connection Between the Two Tracks

4.1 Same Algorithm Family

PuLP delegates to COIN-CBC, which is a branch-and-cut solver, a Branch & Bound spine extended with cutting planes and heuristics. The driver-to-region ILP is small enough that CBC solves it in milliseconds, but on production-sized instances the same algorithmic structure governs solve time. The knapsack track is therefore not a side study: it dissects the engine that the production track depends on.

The knapsack problem also serves as the right testbed for a future LLM-generated-code study: small, well-defined, with a hand-verifiable ground truth (PLAN-v1 instance, optimal value = 11) and now an automated test suite that detects regressions.

4.2 Algorithm vs. Hardware

The two effects observed in Section 3 (algorithmic speedup and language speedup) can be compared directly by placing all four implementations on a single chart.

[image]

Figure 4. Algorithm beats hardware. PuLP/CBC, written in Python and calling a C solver, finishes well under a second at $n = 25$, approximately $129 \times$ faster than the C++ brute force on the same instance and $580 \times$ faster than the x86-64 ASM brute force. A smarter algorithm in a slower language beats brute force in the fastest one.

Algorithmic choice and language choice are not interchangeable. A constant-factor language win (Section 3.5, plateauing around $21 \times$ for Python vs ASM) is dwarfed by the algorithmic gap (six orders of magnitude at $n = 25$). Both effects are real and both matter, but on exponential problems the algorithm dominates.

5. Verification

A pytest suite (`tests/test_knapsack.py`, **9 tests**) cross-checks four independent implementations of the same problem and validates every returned solution for feasibility. Four angles of verification:

Known-answer tests. PuLP and B&B both return the documented optima on two hand-verifiable instances:

- PLAN-v1 small instance (`weights=[4, 3, 2]`, `values=[8, 5, 6]`, `capacity=5`): `value = 11`
- Disaster-relief instance (10 items, `capacity = 50`): `value = 598`

Cross-solver agreement. On the canonical PLAN-v1 instance, PuLP (CBC), the hand-built B&B, and the brute-force enumerator all return `value = 11`. Each reported solution is then independently re-checked against the knapsack invariants.

Random-instance fuzzing. 50 seeded random instances (`random.Random(42)`, $n \in [1, 12]$, integer weights in $[1, 20]$, integer values in $[1, 100]$, `capacity` in $[1, \Sigma \text{weights}]$) are solved by both B&B and brute force; the two must return the same optimal value on every instance. Brute force is the **oracle** here. It is correct by exhaustive enumeration, so any divergence is a B&B bug.

Solution invariants. A shared `_assert_valid_solution` helper checks every reported solution for: (a) distinct indices, (b) total weight $\leq$ capacity, (c) reported value = $\Sigma(\text{values of selected items})$. This catches a class of bugs that a value-only check would miss. For example, a solver that reports the right optimum but the wrong item list.

```
$ pytest -v
tests/test_knapsack.py::test_pulp_plan_v1_instance
tests/test_knapsack.py::test_pulp_disaster_relief_instance
tests/test_knapsack.py::test_bb_matches_pulp_on_plan_v1
tests/test_knapsack.py::test_bb_matches_pulp_on_disaster_relief
tests/test_knapsack.py::test_reader_accepts_three_line_format
tests/test_knapsack.py::test_reader_accepts_four_line_format
tests/test_knapsack.py::test_reader_rejects_malformed
tests/test_knapsack.py::test_all_solvers_agree_on_canonical_instance
tests/test_knapsack.py::test_bb_matches_brute_force_on_random_instances
```

===== 9 passed in 0.23s =====

The full suite runs in under a quarter of a second and is part of `make test`, so any change to a solver re-runs all 9 checks automatically.

6. Reproducibility

Command	Effect
<code>make setup</code>	install dependencies via Poetry
<code>make test</code>	run the pytest suite
<code>python src/main.py driver-region</code>	solve the ILP on the sample CSV
<code>python src/main.py knapsack-pulp</code>	solve PLAN-v1 and disaster-relief via PuLP
<code>python src/main.py knapsack-bb --input data/knapsack_input.txt</code>	solve via hand-built B&B

7. Limitations & Future Work

- **LLM-generated-solver comparison** was the originally proposed Phase 4 of this project and was not completed within this iteration. The hand-built B&B suite serves as ground truth for a future study that asks Claude (or another LLM) to generate a solver and measures whether it finds B&B or defaults to brute force. The verification harness in Section 5 was built specifically so this comparison can be automated.
- **Single-run benchmarks.** All times reported in Section 3.3 are single-run measurements. A more rigorous protocol would re-run each instance 10–30 times and report the median with an inter-quartile range, removing the small per-run noise visible at the timer-floor end of the data.
- **Larger instance sizes.** Brute force was capped at $n = 25$ because the curve makes the point and the larger n add no information. B&B finishes in microseconds at $n = 25$ and could be pushed to $n = 100+$ to characterize the long-tail scaling of the pruning strategy itself, which would be useful for the future LLM comparison.
- **Adversarial knapsack families.** Strongly correlated, inverse strongly correlated, and subset-sum instances (Pisinger 1997) are designed to defeat LP-relaxation pruning. The current synthetic instances are uncorrelated and easy to prune.
- **OR-Library benchmark instances** were planned but not run; synthetic $n = 10\dots 25$ instances were sufficient to demonstrate the algorithmic effect.
- **Hexaly** could not be installed locally and was replaced by the PuLP/CBC track. The originally written `src/knapsack.hxm` is preserved for future reference.
- The driver-region model uses a synthetic 4×3 dataset; scaling to a realistic dataset would stress CBC on the same problem structure and produce a directly comparable B&B curve to the knapsack track.

8. Conclusion

My main result is a verified $\sim 10^6 \times$ speedup on $n = 25$ from B&B over brute force, holding both the language and the machine constant. Language gains exist but are dwarfed by the algorithmic gain. Algorithmic choice dominates language choice. The same B&B algorithm sits inside CBC, the solver that handles the driver-to-region track. The two tracks of this

capstone are not separate projects but two views of the same machinery.

Apart from that main number, the report also contributes a small but defensible verification harness: nine pytest checks that cross-validate four independent solvers (PuLP, brute force, hand-built B&B, plus the format reader) and assert per-solution feasibility on every reported result. This harness is what makes the future LLM-generated-solver comparison (Section 7) a measurable experiment rather than an anecdote.

9. References

1. Land, A. H., and Doig, A. G. (1960). *An automatic method for solving discrete programming problems*. *Econometrica* 28(3), 497–520. Original Branch & Bound paper.
2. Dantzig, G. B. (1957). *Discrete-variable extremum problems*. *Operations Research* 5(2), 266–277. LP relaxation as a bound for integer programs; the fractional-knapsack closed form used here.
3. Kellerer, H., Pferschy, U., and Pisinger, D. (2004). *Knapsack Problems*. Springer. Canonical textbook covering brute force, dynamic programming, branch-and-bound, and benchmark families.
4. Pisinger, D. (1997). *A minimal algorithm for the 0-1 knapsack problem*. *Operations Research* 45(5), 758–767. `minknap`; state-of-the-art for hard knapsack instances.
5. Forrest, J. J., and Lougee-Heimer, R. (2005). *CBC user guide*. INFORMS Tutorials in Operations Research. The branch-and-cut solver that PuLP delegates to in this project.
6. Mitchell, S., O'Sullivan, M., and Dunning, I. (2011). *PuLP: A linear programming toolkit for Python*. The modeling layer used in the production track.